

Implementing matrix multiplication using several approaches and languages.

### **Single-Threaded Implementations:**

#### **C (matrix\_multi\_no\_threads.c):**

Implements matrix multiplication with a standard three-loop algorithm. Performs multiplication with random generated matrices, and measures execution time for increasing matrix sizes.

#### **Python (matrix\_multi\_no\_threads.py):**

Nested loops to multiply matrices in Python. A small  $3 \times 3$  case using matrices filled with ones is used for correctness verification, and then larger random matrices are multiplied while timing the execution.

### **Multithreaded Implementations:**

#### **C with OpenMP (matrix\_multi\_openmp.c):**

OpenMP is used to parallelize the matrix multiplication loops. The number of threads adjusts dynamically, and results are logged for different thread counts.

#### **C with Pthreads (matrix\_multi\_pthreads.c):**

Partitions the work by dividing the matrix rows among threads

#### **Python with Threads (matrix\_multi\_threads.py):**


Uses Python's `concurrent.futures.ThreadPoolExecutor` to split the multiplication task among multiple threads

### **Automation**

**run\_tests\_og.sh** script compiles and runs the C programs and executes the Python programs. Aggregates the results into a log file.

The **timing\_comparison.log** file contains the detailed timing results for all tests, covering a range of matrix sizes (from  $3 \times 3$  to  $1500 \times 1500$ ) and varying thread counts.


Images:



- **Blue (C No Threads):** Single-threaded C implementation. Execution time grows significantly once the matrix size goes 400
- **Orange (C PThreads):** Using the PThreads library for multithreading, stays low even as the matrix size increases, demonstrating how parallelizing the workload can reduce time.
- **Green (C OpenMP):** Similar to PThreads, OpenMP also provides a parallelized approach. Its performance is close to PThreads for most sizes, outperforming both the single-threaded version as the matrix size grows.

### Conclusion:

Multithreaded versions (PThreads and OpenMP) scale better for large matrices compared to the single-threaded C code, resulting in lower execution times.



- **Blue (Python No Threads):** Shows the timing for Python's, single-threaded implementation. It is very fast for small matrices but execution time rises quickly as the matrix size grows. Matrix 500×1500, it reaches over 100 seconds.
- **Orange (Python Threads):** Uses Python's threading (`ThreadPoolExecutor`). A clear improvement over the single-threaded version, the execution times are still much higher than the C versions

### Conclusion:

Although threading helps reduce the time in Python, Python loops limit its performance gains compared to C. These two figures highlight the strong advantage of multithreading in C and demonstrate that Python can benefit from threading; it is still behind C's performance for numerical computations.

## Conclusion

- **Single-Threaded Implementations:**

The C implementation (matrix\_multi\_no\_threads.c) is efficient, even for larger matrix sizes, while the Python implementation (matrix\_multi\_no\_threads.py) suffers as the size increases.

- **Multithreaded Implementations:**

The Pthreads-based C implementation when multiple threads are used improves execution time. But there is an optimal range for thread count, too many threads can limit performance.

- **General Observations:**

OpenMP or Pthreads provides clear performance advantages, python offers rapid development, but its loop-based implementations are not best for high-performance computing unless optimized libraries ( NumPy) are used. Parallelism is key to scaling up the performance.